

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
In [ ]: from google.colab import drive
drive.mount('/content/drive', force_remount=True)
```

Mounted at /content/drive

```
In [ ]: # packages
import numpy as np
import pandas as pd
import os
import json
import seaborn as sns
import matplotlib.pyplot as plt
import re
```

```
In [ ]: # data
file = "/content/drive/MyDrive/movie_dataset_public_final"
```

```
In [ ]: base_dir = "/content/drive/MyDrive/movie_dataset_public_final/raw/"
```

```
def load_ndjson_file(file_name, limit=None):
    data = []
    file_path = os.path.join(base_dir, file_name)

    with open(file_path, 'r', encoding='utf-8') as f:
        for i, line in enumerate(f):
            if limit and i >= limit:
                break
            try:
                data.append(json.loads(line))
            except json.JSONDecodeError as e:
                print(f"Error parsing line {i}: {line}")
                print(f"Error: {e}")
    return data

metadata = load_ndjson_file("metadata.json")
```

```
In [ ]: ratings = load_ndjson_file("ratings.json", limit = 5000000)
print(ratings[:3])
ratings = pd.DataFrame(ratings)
ratings.head()
```

```
[{'item_id': 5, 'user_id': 997206, 'rating': 3.0}, {'item_id': 10, 'user_id': 997206, 'rating': 4.0}, {'item_id': 13, 'user_id': 997206, 'rating': 4.0}]
```

```
Out[ ]: item_id user_id rating
```

0	5	997206	3.0
1	10	997206	4.0
2	13	997206	4.0
3	17	997206	5.0
4	21	997206	4.0

select users >= 100 and movies >= 70

```
In [ ]: user_counts = ratings['user_id'].value_counts()
valid_users = user_counts[user_counts > 100].index
qualified_ratings = ratings[ratings['user_id'].isin(valid_users)]
```

```
In [ ]: movie_counts = qualified_ratings['item_id'].value_counts()
valid_movies = movie_counts[movie_counts > 70].index
qualified_ratings = qualified_ratings[qualified_ratings['item_id'].isin(valid_movies)]
```

```
In [ ]: qualified_ratings
```

```
Out [ ]:
```

	item_id	user_id	rating
0	5	997206	3.0
1	10	997206	4.0
2	13	997206	4.0
3	17	997206	5.0
4	21	997206	4.0
...
4999944	3444	674044	3.0
4999945	3448	674044	4.0
4999946	3519	674044	4.0
4999947	4006	674044	4.0
4999948	5060	674044	4.0

3001682 rows × 3 columns

split into test and train dataset

```
In [ ]: from sklearn.model_selection import train_test_split
train_data, test_data = train_test_split(qualified_ratings, test_size=0.2, random_state=10)
```

```
In [ ]: train_file_path = '/content/drive/My Drive/movie_dataset_public_final/train_data_new.csv'
test_file_path = '/content/drive/My Drive/movie_dataset_public_final/test_data_new.csv'

train_data.to_csv(train_file_path, index=False)
test_data.to_csv(test_file_path, index=False)

print(f"Train data saved to Google Drive at: {train_file_path}")
print(f"Test data saved to Google Drive at: {test_file_path}")
```

Train data saved to Google Drive at: /content/drive/My Drive/movie_dataset_public_final/train_data_new.csv
Test data saved to Google Drive at: /content/drive/My Drive/movie_dataset_public_final/test_data_new.csv

```
In [ ]: train_data
```

```
Out [ ]:
```

	item_id	user_id	rating
2698127	4386	288742	1.0
4729044	1466	209015	3.0
4567052	1446	29651	4.0
3360567	3307	711179	5.0
3574627	6639	89881	4.5
...
2806592	3730	861222	3.0
3609440	11	270526	3.0
4627188	2135	529638	3.0
3549771	898	699790	5.0
2380315	728	451757	5.0

2401345 rows × 3 columns

```
In [ ]: user_movie_matrix = train_data.pivot_table(index='user_id', columns='item_id', values='rating')
user_movie_matrix
```

```
Out [ ]:
```

item_id	1	2	3	4	5	6	7	8	9	10	...	122882	122886	122904	134130	134853	142488	14
user_id																		
19	5.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	3.0	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN
110	4.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	3.0	3.0	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN
144	3.0	NaN	NaN	NaN	NaN	5.0	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN
194	NaN	4.0	NaN	NaN	4.0	NaN	4.0	NaN	NaN	3.0	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN
281	NaN	3.0	NaN	NaN	NaN	NaN	NaN	NaN	4.0	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN
...
999799	5.0	NaN	NaN	NaN	NaN	3.0	NaN	NaN	NaN	4.0	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN
999845	NaN	NaN	NaN	NaN	NaN	3.0	1.0	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN
999851	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN
999869	NaN	NaN	NaN	NaN	NaN	3.0	2.0	NaN	NaN	2.0	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN
999901	5.0	NaN	NaN	NaN	NaN	4.0	3.0	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN

12563 rows × 3804 columns

```
In [ ]:
```

```
# Normalize Ratings by Subtracting User Means
user_means = user_movie_matrix.mean(axis=1) # Calculate mean rating for each user
user_movie_matrix = user_movie_matrix.sub(user_means, axis=0) # Subtract user means
user_movie_matrix.head()
```

```
Out [ ]:
```

item_id	1	2	3	4	5	6	7	8	9	10	...	122882	122886	1229
user_id														
19	1.406114	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	-0.593886	...	NaN	NaN	N
110	0.872247	NaN	NaN	NaN	NaN	NaN	NaN	NaN	-0.127753	-0.127753	...	NaN	NaN	N
144	-0.174212	NaN	NaN	NaN	NaN	1.825788	NaN	NaN	NaN	NaN	...	NaN	NaN	N
194	NaN	0.285714	NaN	NaN	0.285714	NaN	0.285714	NaN	NaN	-0.714286	...	NaN	NaN	N
281	NaN	-0.816667	NaN	NaN	NaN	NaN	NaN	NaN	0.183333	NaN	...	NaN	NaN	N

5 rows × 3804 columns

```
In [ ]:
```

```
# Check sparsity
num_total_entries = user_movie_matrix.size # Total number of cells
num_missing_entries = user_movie_matrix.isna().sum().sum() # Number of missing entries
sparsity = num_missing_entries / num_total_entries # Sparsity calculation

print(f"Data sparsity: {sparsity:.2%}") # Output sparsity as a percentage
```

Data sparsity: 95.00%

use svd to fulfill NAN values within user-item matrix

```
In [ ]:
```

```
user_movie_matrix.shape
```

```
Out [ ]:
```

```
(12563, 3804)
```

```
In [ ]:
```

```
from sklearn.decomposition import TruncatedSVD
import numpy as np
from tqdm import tqdm
import time

# Fill NaN values in the user-item matrix with 0 (as required for SVD)
user_movie_matrix_filled = user_movie_matrix.fillna(0).values

def svd_with_progress(data, n_components, chunk_size=1000):
    """
    Perform SVD with progress tracking using tqdm.

    Parameters:
    - data: The user-item matrix (numpy array).
    - n_components: Number of components for dimensionality reduction.
    - chunk_size: Number of rows to process at a time.
```

```

Returns:
- user_features: User embedding matrix.
- movie_features: Movie embedding matrix (transposed).
"""

svd = TruncatedSVD(n_components=n_components, random_state=42)
user_features_list = []
start_time = time.time()

# Process the matrix in chunks
for i in tqdm(range(0, data.shape[0], chunk_size), desc="Processing SVD"):
    chunk = data[i:i + chunk_size]
    user_features_list.append(svd.fit_transform(chunk))

# Combine processed chunks
user_features = np.vstack(user_features_list)
end_time = time.time()

print(f"SVD completed in {end_time - start_time:.2f} seconds")
return user_features, svd.components_.T

# Execute SVD decomposition
n_components = 50 # Adjust based on your dataset
user_features, movie_features = svd_with_progress(
    user_movie_matrix_filled, n_components=n_components, chunk_size=1000
)

# Predict the rating matrix by multiplying user and movie features
predicted_ratings = np.dot(user_features, movie_features.T)

# Convert predicted ratings back to a DataFrame for easier interpretation
predicted_ratings_df = pd.DataFrame(predicted_ratings, index=user_movie_matrix.index, columns=user_movie_matrix.columns)

# Display progress
print("Predicted ratings matrix:")
print(predicted_ratings_df.head())

```

Processing SVD: 100%|██████████| 13/13 [00:11<00:00, 1.16it/s]

SVD completed in 11.26 seconds

Predicted ratings matrix:

item_id	1	2	3	4	5	6	7	\
user_id								
19	0.439225	-0.210008	-0.067958	-0.195623	-0.249413	0.249318	0.108626	
110	0.226099	-0.258291	0.022680	-0.139956	-0.178066	0.208525	-0.134432	
144	3.928778	0.676412	0.465336	-0.290208	-0.085016	-0.820204	-0.559409	
194	0.372853	0.064778	0.004215	-0.029979	-0.000832	0.135179	-0.008437	
281	0.126331	-0.025364	0.023498	-0.024513	0.079875	0.101693	0.024951	

item_id	8	9	10	...	122882	122886	122904	\
user_id				...				
19	-0.042914	-0.186523	0.029838	...	0.025267	-0.035646	0.004302	
110	-0.064228	-0.033422	0.080660	...	-0.053851	0.006396	0.074810	
144	-0.285942	-0.237776	-0.236610	...	-0.138309	0.121350	-0.074112	
194	-0.046637	-0.068643	-0.011595	...	-0.032872	0.025768	-0.013631	
281	0.037433	-0.092773	0.198741	...	0.000720	0.030816	-0.006240	

item_id	134130	134853	142488	143385	148626	164179	166528
user_id							
19	-0.013165	0.032932	0.047320	-0.024802	0.017723	-0.034189	0.003195
110	0.058365	0.024062	0.038030	-0.009597	0.014995	0.019629	0.027969
144	0.071783	0.132894	-0.139098	-0.141977	0.047853	0.151878	-0.072762
194	-0.005684	-0.015055	-0.019203	-0.023534	-0.004363	0.016544	-0.004241
281	0.045348	-0.003681	0.011002	0.047445	-0.012434	0.039675	-0.007875

[5 rows x 3804 columns]

use cosine similarity to calculate the similarity metrics

```

In [ ]: from sklearn.metrics.pairwise import cosine_similarity
import pandas as pd

# Compute user-to-user similarity
user_similarity = cosine_similarity(user_movie_matrix_filled) # Input is the filled user-item matrix
user_similarity_df = pd.DataFrame(user_similarity, index=user_movie_matrix.index, columns=user_movie_matrix.index)

print("User Similarity Matrix:")
print(user_similarity_df.head())

```

User Similarity Matrix:

```

user_id    19      110      144      194      281      301      337      \
user_id
19      1.000000  0.055390  0.082475  0.052292  0.008436  0.091234  0.101988
110     0.055390  1.000000 -0.039201  0.059968  0.057699 -0.005846  0.060626
144     0.082475 -0.039201  1.000000 -0.031112 -0.025069  0.021368  0.033775
194     0.052292  0.059968 -0.031112  1.000000  0.136647  0.041423  0.055263
281     0.008436  0.057699 -0.025069  0.136647  1.000000  0.048269 -0.049340

```

```

user_id    364      372      489      ...    999506    999527    999549    \
user_id
19      0.093516  0.127466  0.026647  ...    0.033266  0.060578  0.025270
110     0.087973  0.031983 -0.022259  ...   -0.017718  0.062271  0.033002
144     0.001275  0.034263  0.041980  ...    0.067173  0.039500  0.011099
194     0.162218 -0.034590 -0.019204  ...   -0.003641  0.007233 -0.032069
281     0.126720  0.079098  0.037195  ...   -0.075035 -0.028103  0.011247

```

```

user_id    999590    999721    999799    999845    999851    999869    999901
user_id
19      0.027110  0.062299  0.099324  0.068719  0.054825  0.035108  0.112084
110     0.107345  0.093093  0.040699  0.002348  0.044720  0.009803  0.031383
144    -0.005779 -0.003979  0.050420  0.101318  0.017539  0.178241  0.077019
194     0.275258  0.164485  0.033521  0.005188  0.030821  0.017537  0.051423
281     0.104222  0.002516  0.054108  0.024325  0.035936 -0.005894  0.038060

```

[5 rows x 12563 columns]

In []: user_similarity_df

```

Out [ ]: user_id    19      110      144      194      281      301      337      364      372      489      ...
user_id
19      1.000000  0.055390  0.082475  0.052292  0.008436  0.091234  0.101988  0.093516  0.127466  0.026647  ...
110     0.055390  1.000000 -0.039201  0.059968  0.057699 -0.005846  0.060626  0.087973  0.031983 -0.022259  ...
144     0.082475 -0.039201  1.000000 -0.031112 -0.025069  0.021368  0.033775  0.001275  0.034263  0.041980  ...
194     0.052292  0.059968 -0.031112  1.000000  0.136647  0.041423  0.055263  0.162218 -0.034590 -0.019204  ...
281     0.008436  0.057699 -0.025069  0.136647  1.000000  0.048269 -0.049340  0.126720  0.079098  0.037195  ...
...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...      ...
999799  0.099324  0.040699  0.050420  0.033521  0.054108  0.060124  0.113892  0.032156  0.008993  0.050281  ...
999845  0.068719  0.002348  0.101318  0.005188  0.024325  0.079532  0.042707 -0.010867  0.035076  0.058930  ...
999851  0.054825  0.044720  0.017539  0.030821  0.035936  0.126800  0.016192  0.111933  0.012182  0.000232  ...
999869  0.035108  0.009803  0.178241  0.017537 -0.005894  0.067441  0.057670  0.012516  0.032242  0.072707  ...
999901  0.112084  0.031383  0.077019  0.051423  0.038060  0.092930  0.018396  0.045890  0.095025  0.026879  ...

```

12563 rows x 12563 columns

```

In [ ]: def predict_rating(target_user, target_movie, user_similarity_matrix, user_movie_matrix, K=5):
        """
        Predict the rating of a target_user for a target_movie based on similar users' ratings.
        """
        # Step 1: Get similarity scores for the target user
        similarities = user_similarity_matrix.loc[target_user]

        # Step 2: Extract ratings for the target movie from all users
        movie_ratings = user_movie_matrix[target_movie]

        # Step 3: Filter out users who haven't rated the target movie
        valid_users = movie_ratings[~movie_ratings.isna()].index
        valid_similarities = similarities[valid_users]
        valid_ratings = movie_ratings[valid_users]

        # Step 4: Select top-K similar users
        top_k_users = valid_similarities.nlargest(K)

        # Step 5: Calculate the weighted average rating
        numerator = (top_k_users * valid_ratings[top_k_users.index]).sum()
        denominator = top_k_users.abs().sum()

        # Step 6: Return predicted rating or default value
        if denominator > 0:
            return numerator / denominator

```

```

else:
    return user_movie_matrix.loc[target_user].mean() # Default to the user's mean rating

```

```

In [ ]: # Example inputs
target_user = user_movie_matrix.index[0] # Replace with the desired user ID
target_movie = user_movie_matrix.columns[0] # Replace with the desired movie ID

predicted_rating = predict_rating(target_user, target_movie, user_similarity_df, user_movie_matrix)
print(f"Predicted rating for user {target_user} on movie {target_movie}: {predicted_rating}")

```

Predicted rating for user 19 on movie 1: 1.0133675518654128

```

In [ ]: train_data[qualified_ratings['user_id']==19]

```

<ipython-input-52-6132d34bd993>:1: UserWarning: Boolean Series key will be reindexed to match DataFrame index.
train_data[qualified_ratings['user_id']==19]

```

Out[ ]:
   item_id  user_id  rating
1211634    117      19     3.0
1211852   1276      19     5.0
1211689    357      19     3.0
1211724    524      19     4.0
1211716    480      19     3.0
...      ...      ...     ...
1211705    434      19     4.0
1211713    474      19     3.0
1211621     45      19     3.0
1211722    509      19     3.0
1211718    491      19     4.0

```

229 rows × 3 columns

model evaluation

```

In [ ]: from tqdm import tqdm
import time

def safe_predict_rating(row):
    try:
        target_user = row['user_id']
        target_movie = row['item_id']

        return predict_rating(target_user, target_movie, user_similarity_df, user_movie_matrix)
    except Exception as e:
        print(f"Error predicting for user {row['user_id']} and movie {row['item_id']}: {e}")
        return default_value

start_time = time.time()

default_value = 3.0

tqdm.pandas(desc="Predicting Ratings")
train_data['predicted_rating'] = train_data.progress_apply(safe_predict_rating, axis=1)

end_time = time.time()

print(f"Predicting ratings completed in {end_time - start_time:.2f} seconds")

```

Predicting Ratings: 100%|██████████| 2401345/2401345 [1:55:26<00:00, 346.71it/s]
Predicting ratings completed in 6926.12 seconds

```

In [ ]: train_data

```

```
Out [ ]:
```

	item_id	user_id	rating	predicted_rating	
	2698127	4386	288742	1.0	-1.339137
	4729044	1466	209015	3.0	-0.224275
	4567052	1446	29651	4.0	0.453455
	3360567	3307	711179	5.0	0.948223
	3574627	6639	89881	4.5	0.640596

	2806592	3730	861222	3.0	0.078697
	3609440	11	270526	3.0	-0.421295
	4627188	2135	529638	3.0	-0.449507
	3549771	898	699790	5.0	0.767411
	2380315	728	451757	5.0	0.890322

2401345 rows x 4 columns

```
In [ ]: # Calculate the average rating per user
user_avg_rating = train_data.groupby('user_id')['rating'].mean()

# Map the average rating back to the DataFrame
train_data['average_rating_x'] = train_data['user_id'].map(user_avg_rating)
```

```
In [ ]: #add average rating back
train_data.loc[:, 'restored_predicted_rating'] = (
    train_data['predicted_rating'] + train_data['average_rating_x']
)
```

```
In [ ]: #debug: ValueError: Input contains NaN.
print(train_data[['rating', 'restored_predicted_rating']].isna().sum())
```

```
rating          0
restored_predicted_rating  0
dtype: int64
```

```
In [ ]: file_path = '/content/drive/My Drive/train_data_with_restored_data_new.csv'

train_data.to_csv(file_path, index=False)

print(f"File saved to Google Drive at: {file_path}")
```

File saved to Google Drive at: /content/drive/My Drive/train_data_with_restored_data.csv

```
In [ ]: # mount google drive
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [ ]: from google.colab import drive
drive.mount('/content/drive', force_remount=True)
```

Mounted at /content/drive

```
In [ ]: import pandas as pd

file_path = '/content/drive/MyDrive/train_data_with_restored_data_new.csv'
train_data = pd.read_csv(file_path)
print("Data loaded successfully!")
```

Data loaded successfully!

```
In [ ]: from sklearn.metrics import mean_squared_error, mean_absolute_error
import numpy as np
```

```
In [ ]: cleaned_ratings = train_data.dropna(subset=['restored_predicted_rating'])

y_true = cleaned_ratings['rating']
y_pred = cleaned_ratings['restored_predicted_rating']

mse = mean_squared_error(y_true, y_pred)
rmse = np.sqrt(mse)
mae = mean_absolute_error(y_true, y_pred)

print(f"MSE: {mse:.4f}")
print(f"RMSE: {rmse:.4f}")
print(f"MAE: {mae:.4f}")
```

MSE: 0.1145
RMSE: 0.3384
MAE: 0.2612

```
In [ ]: cleaned_ratings
```

```
Out [ ]:
```

	item_id	user_id	rating	predicted_rating	average_rating_x	restored_predicted_rating
0	4386	288742	1.0	-1.339137	3.266987	1.927850
1	1466	209015	3.0	-0.224275	3.151261	2.926985
2	1446	29651	4.0	0.453455	3.725248	4.178703
3	3307	711179	5.0	0.948223	3.921965	4.870189
4	6639	89881	4.5	0.640596	3.982219	4.622815
...
2401340	3730	861222	3.0	0.078697	3.236501	3.315198
2401341	11	270526	3.0	-0.421295	3.273973	2.852678
2401342	2135	529638	3.0	-0.449507	3.397727	2.948220
2401343	898	699790	5.0	0.767411	4.125000	4.892411
2401344	728	451757	5.0	0.890322	3.497265	4.387588

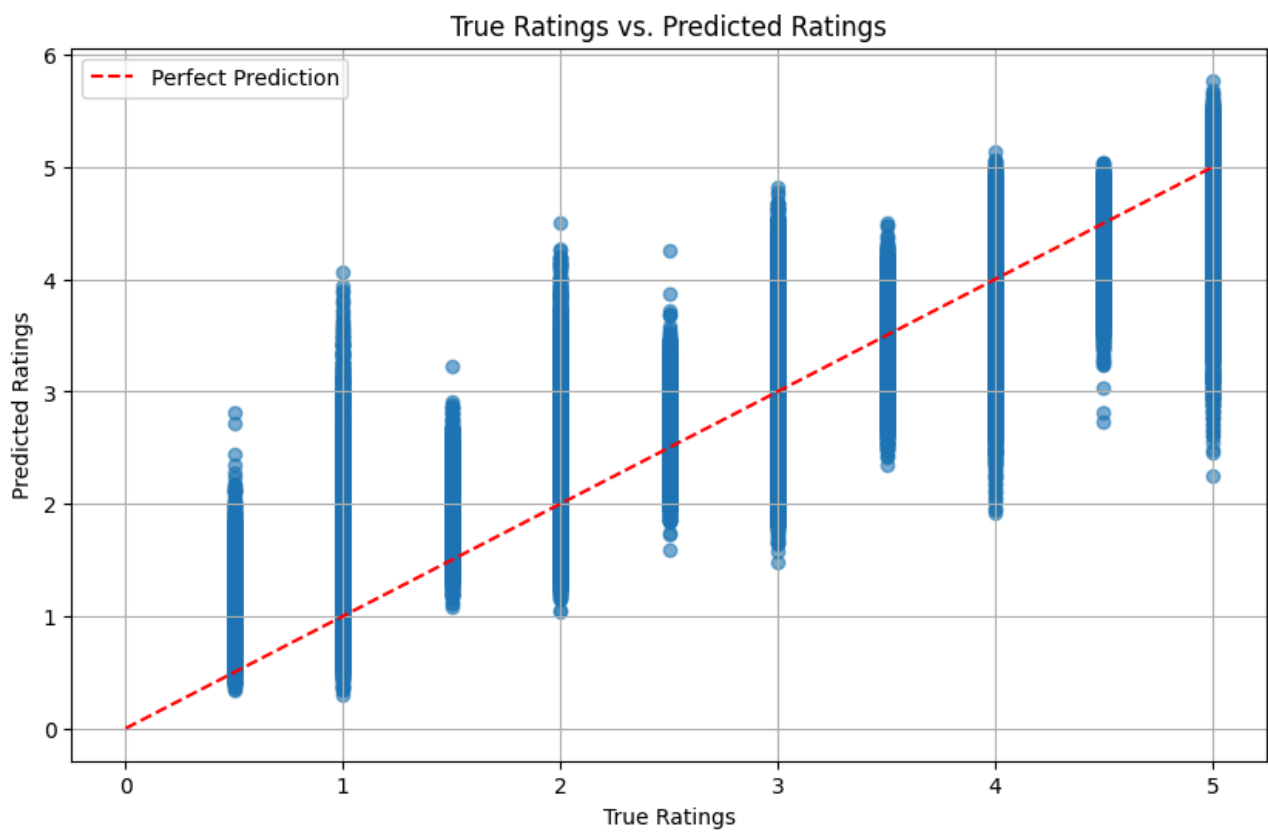
2401345 rows × 6 columns

```
In [ ]: import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))

plt.scatter(y_true, y_pred, alpha=0.6)
plt.plot([0, 5], [0, 5], color='red', linestyle='--', label='Perfect Prediction')

plt.title("True Ratings vs. Predicted Ratings")
plt.xlabel("True Ratings")
plt.ylabel("Predicted Ratings")
plt.legend()
plt.grid()
plt.show()
```

```
In [ ]: # Save the DataFrame as a .pkl file in Google Drive
file_path = '/content/drive/My Drive/cleaned_ratings_cosine_train.pkl'
cleaned_ratings.to_pickle(file_path)
```

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
In [ ]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.preprocessing import OneHotEncoder, LabelEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import matplotlib.pyplot as plt
import seaborn as sns
import os
import json
from tqdm import tqdm
```

```
In [ ]: base_dir = "/content/drive/MyDrive/movie_dataset_public_final/raw/"

def load_ndjson_file(file_name, limit=None):
    data = []
    file_path = os.path.join(base_dir, file_name)

    with open(file_path, 'r', encoding='utf-8') as f:
        for i, line in enumerate(f):
            if limit and i >= limit:
                break
            try:
                data.append(json.loads(line))
            except json.JSONDecodeError as e:
                print(f"Error parsing line {i}: {line}")
                print(f"Error: {e}")

    return data

metadata = load_ndjson_file("metadata.json")
```

```
In [ ]: metadata_updated = load_ndjson_file("metadata_updated.json")
metadata = pd.DataFrame(metadata_updated)
print(metadata.head())
```

	title	directedBy	\
0	Toy Story (1995)	John Lasseter	
1	Jumanji (1995)	Joe Johnston	
2	Grumpier Old Men (1995)	Howard Deutch	
3	Waiting to Exhale (1995)	Forest Whitaker	
4	Father of the Bride Part II (1995)	Charles Shyer	

	starring	avgRating	imdbId
0	Tim Allen, Tom Hanks, Don Rickles, Jim Varney,...	3.89146	0114709
1	Jonathan Hyde, Bradley Pierce, Robin Williams,...	3.26605	0113497
2	Jack Lemmon, Walter Matthau, Ann-Margret , Sop...	3.17146	0113228
3	Angela Bassett, Loretta Devine, Whitney Housto...	2.86824	0114885
4	Steve Martin, Martin Short, Diane Keaton, Kimb...	3.07620	0113041

	item_id
0	1
1	2
2	3
3	4
4	5

```
In [ ]: metadata['release_year'] = metadata['title'].str.extract(r'\((\d{4})\)'),
metadata['movie_title'] = metadata['title'].str.replace(r'\((\d{4})\)',' ',
metadata = metadata.drop(columns=['title'])

print(metadata.head())
```

	directedBy	starring	\
0	John Lasseter	Tim Allen, Tom Hanks, Don Rickles, Jim Varney,...	
1	Joe Johnston	Jonathan Hyde, Bradley Pierce, Robin Williams,...	
2	Howard Deutch	Jack Lemmon, Walter Matthau, Ann-Margret , Sop...	
3	Forest Whitaker	Angela Bassett, Loretta Devine, Whitney Housto...	
4	Charles Shyer	Steve Martin, Martin Short, Diane Keaton, Kimb...	

	avgRating	imdbId	item_id	release_year	movie_title
0	3.89146	0114709	1	1995	Toy Story
1	3.26605	0113497	2	1995	Jumanji
2	3.17146	0113228	3	1995	Grumpier Old Men
3	2.86824	0114885	4	1995	Waiting to Exhale
4	3.07620	0113041	5	1995	Father of the Bride Part II

```
In [ ]: # Handle missing values
metadata['directedBy'].fillna('Unknown', inplace=True)

metadata['directedBy'] = metadata['directedBy'].str.strip()
metadata['directedBy'] = metadata['directedBy'].str.replace(r"\s+", " ",

metadata['directedBy'] = metadata['directedBy'].str.replace("&", ",") #
metadata['directedBy'] = metadata['directedBy'].str.split(",").apply(
    lambda x: [d.strip() for d in x if d.strip()] if isinstance(x, list)
)

metadata['directedBy'] = metadata['directedBy'].apply(lambda x: ", ".join
```

<ipython-input-6-2ec3f22f36bd>:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
metadata['directedBy'].fillna('Unknown', inplace=True)
```

```
In [ ]: tag_count = load_ndjson_file("tag_count.json")
tags = load_ndjson_file("tags.json")
```

```
In [ ]: tags = pd.DataFrame(tags)
tag_count = pd.DataFrame(tag_count)

tags_joined = pd.merge(tags, tag_count, left_on='id', right_on='tag_id',
tags_joined = tags_joined.drop(columns=['id'])

print(tags_joined.head())
```

	tag	item_id	tag_id	num
0	aardman	720	22	5
1	aardman	745	22	22
2	aardman	1148	22	19
3	aardman	1223	22	9
4	aardman	3429	22	3

```
In [ ]: tags_grouped = tags_joined.groupby('item_id')['tag'].apply(list).reset_index()
print(tags_grouped.head())
```

	item_id	tag
0	1	[nostalgic, interesting, children, witty, emot...
1	2	[children, animals, based on a book, comedy, a...
2	3	[sequel, good soundtrack, comedy, funny, funni...
3	4	[divorce, revenge, chick flick]
4	5	[remake, touching, sequel, comedy, pregnancy, ...]

```
In [ ]: sentiments_df = pd.read_csv('/content/drive/MyDrive/movie_dataset_public_
```

```
merged_data = pd.merge(metadata, sentiments_df, on='item_id', how='inner')
merged_df = pd.merge(merged_data, tags_grouped, on='item_id', how='inner')
print(merged_df.head())
```

	directedBy	starring \
0	John Lasseter	Tim Allen, Tom Hanks, Don Rickles, Jim Varney,...
1	Joe Johnston	Jonathan Hyde, Bradley Pierce, Robin Williams,...
2	Howard Deutch	Jack Lemmon, Walter Matthau, Ann-Margret , Sop...
3	Forest Whitaker	Angela Bassett, Loretta Devine, Whitney Housto...
4	Charles Shyer	Steve Martin, Martin Short, Diane Keaton, Kimb...

	avgRating	imdbId	item_id	release_year	movie_title
0	3.89146	0114709	1	1995	Toy Story
1	3.26605	0113497	2	1995	Jumanji
2	3.17146	0113228	3	1995	Grumpier Old Men
3	2.86824	0114885	4	1995	Waiting to Exhale
4	3.07620	0113041	5	1995	Father of the Bride Part II

	avg_negative	avg_neutral	avg_positive	avg_compound	\
0	0.042831	0.757843	0.199306	0.904176	
1	0.064142	0.746503	0.189320	0.799850	
2	0.066300	0.744933	0.188833	0.835292	
3	0.080184	0.751204	0.168429	0.513924	
4	0.049833	0.740121	0.210000	0.825379	

	tag
0	[nostalgic, interesting, children, witty, emot...
1	[children, animals, based on a book, comedy, a...
2	[sequel, good soundtrack, comedy, funny, funni...
3	[divorce, revenge, chick flick]
4	[remake, touching, sequel, comedy, pregnancy, ...

```
In [ ]: file_path = '/content/drive/MyDrive/movie_dataset_public_final/qualified
try:
    df = pd.read_csv(file_path)
    print("File imported successfully.")
    # Now you can work with the DataFrame 'df'
except FileNotFoundError:
    print(f"Error: File not found at {file_path}")
except Exception as e:
    print(f"An error occurred: {e}")
```

File imported successfully.

```
In [ ]: movies = df.item_id.unique()
```

```
In [ ]: filtered_merged_df = merged_df[merged_df['item_id'].isin(movies)]
print(filtered_merged_df.head())
```

	directedBy	starring \
0	John Lasseter	Tim Allen, Tom Hanks, Don Rickles, Jim Varney,...
1	Joe Johnston	Jonathan Hyde, Bradley Pierce, Robin Williams,...
2	Howard Deutch	Jack Lemmon, Walter Matthau, Ann-Margret , Sop...
3	Forest Whitaker	Angela Bassett, Loretta Devine, Whitney Housto...
4	Charles Shyer	Steve Martin, Martin Short, Diane Keaton, Kimb...

	avgRating	imdbId	item_id	release_year	movie_title
0	3.89146	0114709	1	1995	Toy Story
1	3.26605	0113497	2	1995	Jumanji
2	3.17146	0113228	3	1995	Grumpier Old Men
3	2.86824	0114885	4	1995	Waiting to Exhale
4	3.07620	0113041	5	1995	Father of the Bride Part II

	avg_negative	avg_neutral	avg_positive	avg_compound	\
0	0.042831	0.757843	0.199306	0.904176	
1	0.064142	0.746503	0.189320	0.799850	
2	0.066300	0.744933	0.188833	0.835292	
3	0.080184	0.751204	0.168429	0.513924	
4	0.049833	0.740121	0.210000	0.825379	

	tag
0	[nostalgic, interesting, children, witty, emot...
1	[children, animals, based on a book, comedy, a...
2	[sequel, good soundtrack, comedy, funny, funni...
3	[divorce, revenge, chick flick]
4	[remake, touching, sequel, comedy, pregnancy, ...

```
In [ ]: from sklearn.preprocessing import OneHotEncoder
import pandas as pd

# Limit to top 50 directors
top_directors = filtered_merged_df['directedBy'].value_counts().nlargest(50)
filtered_merged_df['directedBy'] = filtered_merged_df['directedBy'].apply(lambda x: top_directors.get(x, ''))

director_encoder = OneHotEncoder(handle_unknown='ignore', sparse_output=False)
directors_encoded = director_encoder.fit_transform(filtered_merged_df[['directedBy']])
director_feature_names = director_encoder.get_feature_names_out(['directedBy'])

print(f"Number of encoded features: {directors_encoded.shape[1]}")
print("Feature names:", director_feature_names[:5], "...")

directors_df = pd.DataFrame(directors_encoded, columns=director_feature_names)

filtered_merged_df = pd.concat([filtered_merged_df, directors_df], axis=1)
filtered_merged_df.drop('directedBy', axis=1, inplace=True)
```

```
Number of encoded features: 51
Feature names: ['directedBy_Alan Parker' 'directedBy_Alfred Hitchcock'
'directedBy_Ang Lee' 'directedBy_Barry Levinson'
'directedBy_Billy Wilder'] ...
```

```
In [ ]: filtered_merged_df.shape
```

```
Out[ ]: (3803, 62)
```

```
In [ ]: from sklearn.preprocessing import MultiLabelBinarizer

filtered_merged_df['starring'] = filtered_merged_df['starring'].apply(lambda x: MultiLabelBinarizer().fit_transform(x))
```

```

# Limit to top 50 actors
all_actors = filtered_merged_df.explode('starring')['starring']
top_actors = all_actors.value_counts().nlargest(50).index
filtered_merged_df['starring'] = filtered_merged_df['starring'].apply(lambda x: x if x in top_actors else None)

mlb = MultiLabelBinarizer()
actors_encoded = mlb.fit_transform(filtered_merged_df['starring'])
actor_feature_names = mlb.classes_
actors_df = pd.DataFrame(actors_encoded, columns=actor_feature_names, index=filtered_merged_df.index)

filtered_merged_df = pd.concat([filtered_merged_df, actors_df], axis=1)
filtered_merged_df.drop('starring', axis=1, inplace=True)

```

```
In [ ]: filtered_merged_df['tag'] = filtered_merged_df['tag'].apply(lambda x: x if x in top_tags else None)
```

```

# Step 1: Limit to Top 100 Tags
all_tags = filtered_merged_df.explode('tag')['tag']
top_tags = all_tags.value_counts().nlargest(100).index
filtered_merged_df['tag'] = filtered_merged_df['tag'].apply(lambda x: x if x in top_tags else None)

mlb = MultiLabelBinarizer()
tags_encoded = mlb.fit_transform(filtered_merged_df['tag'])
tag_feature_names = mlb.classes_

tags_df = pd.DataFrame(tags_encoded, columns=tag_feature_names, index=filtered_merged_df.index)

filtered_merged_df = pd.concat([filtered_merged_df, tags_df], axis=1)
filtered_merged_df.drop('tag', axis=1, inplace=True)

```

```
In [ ]: from sklearn.feature_extraction.text import TfidfVectorizer

tfidf = TfidfVectorizer(max_features=100)

title_tfidf = tfidf.fit_transform(filtered_merged_df['movie_title'])

tfidf_df = pd.DataFrame(title_tfidf.toarray(), columns=tfidf.get_feature_names(), index=filtered_merged_df.index)

filtered_merged_df = pd.concat([filtered_merged_df, tfidf_df], axis=1)
filtered_merged_df.drop('movie_title', axis=1, inplace=True)

```

```
In [ ]: filtered_merged_df['release_year'] = pd.to_numeric(filtered_merged_df['release_year'], errors='coerce')
filtered_merged_df['release_year'].fillna(0, inplace=True)

filtered_merged_df['release_year'] = filtered_merged_df['release_year'].astype(int)

```

<ipython-input-19-f1e9d019446f>:6: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
filtered_merged_df['release_year'].fillna(0, inplace=True)
```

```
In [ ]: filtered_merged_df['release_decade'] = (filtered_merged_df['release_year']

# One-Hot Encode release decade
decade_encoder = OneHotEncoder(handle_unknown='ignore', sparse_output=False)
decades_encoded = decade_encoder.fit_transform(filtered_merged_df[['release_year']])
decade_feature_names = decade_encoder.get_feature_names_out(['release_year'])
decades_df = pd.DataFrame(decades_encoded, columns=decade_feature_names,

# Concatenate and drop original columns
filtered_merged_df = pd.concat([filtered_merged_df, decades_df], axis=1)
filtered_merged_df.drop(['release_year', 'release_decade'], axis=1, inplace=True)
```

```
In [ ]: from sklearn.preprocessing import StandardScaler

numerical_features = ['avg_negative', 'avg_neutral', 'avg_positive', 'avg_rating']

scaler = StandardScaler()

filtered_merged_df[numerical_features] = scaler.fit_transform(filtered_merged_df[numerical_features])
```

```
In [ ]: train_data = pd.read_csv('/content/drive/MyDrive/movie_dataset_public_final.csv')
print(train_data.head())
```

	item_id	user_id	rating
0	4386	288742	1.0
1	1466	209015	3.0
2	1446	29651	4.0
3	3307	711179	5.0
4	6639	89881	4.5

```
In [ ]: cb_training_data = pd.merge(train_data, filtered_merged_df, on='item_id',
print(cb_training_data.head())
```


	item_id	user_id	rating	avgRating	imdbId	avg_negative	avg_neutral
0	4386	288742	1.0	2.55868	0239395	0.443117	-0.832400
1	1466	209015	3.0	3.80767	0119008	-0.816323	0.713760
2	1446	29651	4.0	4.00348	0116790	-1.177937	0.225011
3	3307	711179	5.0	4.10615	0021749	-1.202262	-0.696967
4	6639	89881	4.5	3.92525	0062467	1.024506	-0.291457

	avg_positive	avg_compound	directedBy_Alan Parker	...	\
0	0.314369	-0.409147		0.0	...
1	0.102293	0.648067		0.0	...
2	0.819075	1.036453		0.0	...
3	1.609711	1.298728		0.0	...
4	-0.636241	-0.754227		0.0	...

	release_decade_1920	release_decade_1930	release_decade_1940	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	1.0	0.0	
4	0.0	0.0	0.0	

	release_decade_1950	release_decade_1960	release_decade_1970	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	1.0	0.0	

	release_decade_1980	release_decade_1990	release_decade_2000	\
0	0.0	0.0	1.0	
1	0.0	1.0	0.0	
2	0.0	1.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	release_decade_2010
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0

[5 rows x 322 columns]

```
In [ ]: print(cb_training_data.columns)
```

```
Index(['item_id', 'user_id', 'rating', 'avgRating', 'imdbId', 'avg_negative',
      'avg_neutral', 'avg_positive', 'avg_compound', 'directedBy_Alan Parker',
      ...,
      'release_decade_1920', 'release_decade_1930', 'release_decade_1940',
      'release_decade_1950', 'release_decade_1960', 'release_decade_1970',
      'release_decade_1980', 'release_decade_1990', 'release_decade_2000',
      'release_decade_2010'],
      dtype='object', length=322)
```

```
In [ ]: duplicate_columns = cb_training_data.columns[cb_training_data.columns.duplicated()]
print(f"Duplicate columns: {duplicate_columns}")
```

```
Duplicate columns: Index(['Other', 'blood', 'children', 'death', 'love', 'story', 'war'], dtype='object')
```

```
In [ ]: cb_training_data = cb_training_data.loc[:, ~cb_training_data.columns.duplicated()]
```

```
In [ ]: assert cb_training_data.columns.is_unique, "Duplicate columns still exist"
```

```
In [ ]: cb_training_data.columns = cb_training_data.columns.str.replace('[^A-Za-z]', '')
```

```
In [ ]: assert cb_training_data.columns.is_unique, "Duplicate column names still exist"
```

```
In [ ]: import joblib
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, mean_squared_error
import xgboost as xgb
from tqdm import tqdm
import numpy as np

X = cb_training_data.drop(columns=['imdbId', 'rating', 'avgRating', 'item_id'])
y = cb_training_data['rating']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

xgb_model = xgb.XGBRegressor(
    n_estimators=2000,
    learning_rate=0.05,
    tree_method='hist',
    device='cuda',
    random_state=42
)

print("Training XGBoost model...")

with tqdm(total=xgb_model.n_estimators, desc="Training Progress", unit="estimator"):
    def update_progress(env):
        """Update tqdm progress bar."""
        pbar.update(1)

    xgb_model.fit(
        X_train,
        y_train,
        eval_set=[(X_test, y_test)],
        verbose=False,
    )

# Predict and evaluate
y_pred = xgb_model.predict(X_test)
mae = mean_absolute_error(y_test, y_pred)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))

print(f"CB Model - XGBoost:")
print(f"MAE: {mae:.4f}, RMSE: {rmse:.4f}")

# Save the model
```

```
joblib.dump(xgb_model, '/content/drive/MyDrive/movie_dataset_public_final')
print("XGBoost model saved as xgb_cb_model.pkl")
```

Training XGBoost model...

Training Progress: 0% | 0/2000 [02:19<?, ?iteration/s]
 /usr/local/lib/python3.10/dist-packages/xgboost/core.py:158: UserWarning:
 [04:55:55] WARNING: /workspace/src/common/error_msg.cc:58: Falling back to
 prediction using DMatrix due to mismatched devices. This might lead to high
 memory usage and slower performance. XGBoost is running on: cuda:0, while
 the input data is on: cpu.

Potential solutions:

- Use a data structure that matches the device ordinal in the booster.
- Set the device for booster before call to inplace_predict.

This warning will only be shown once.

```
warnings.warn(msg, UserWarning)
```

CB Model - XGBoost:

MAE: 0.7443, RMSE: 0.9422

XGBoost model saved as xgb_cb_model.pkl

```
In [ ]: from sklearn.model_selection import GridSearchCV

param_grid = {
    'n_estimators': [2000],
    'max_depth': [4, 6, 8],
    'learning_rate': [0.01, 0.05, 0.1],
    'subsample': [0.8, 1.0],
    'colsample_bytree': [0.8, 1.0]
}

xgb_model = xgb.XGBRegressor(tree_method='gpu_hist', random_state=42)

grid_search = GridSearchCV(
    estimator=xgb_model,
    param_grid=param_grid,
    scoring='neg_mean_squared_error',
    cv=5,
    verbose=1,
    n_jobs=-1
)

print("Running GridSearchCV...")
grid_search.fit(X_train, y_train)

best_model = grid_search.best_estimator_
print("Best parameters:", grid_search.best_params_)

y_pred = best_model.predict(X_test)
mae = mean_absolute_error(y_test, y_pred)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print(f"Tuned CB Model - XGBoost (CV): MAE: {mae:.4f}, RMSE: {rmse:.4f}")

joblib.dump(best_model, '/content/drive/MyDrive/movie_dataset_public_final')
print("Tuned XGBoost model saved as xgb_cb_model_tuned.pkl")
```

Running GridSearchCV...

Fitting 5 folds for each of 36 candidates, totalling 180 fits

In []:

```
In [ ]: import numpy as np
import pandas as pd
import os
import json
import seaborn as sns
import matplotlib.pyplot as plt
import re
```

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
In [ ]: import os

content_model_path = '/content/drive/MyDrive/movie_dataset_public_final/x
cf_predictions_path = '/content/drive/MyDrive/movie_dataset_public_final/

if os.path.exists(content_model_path):
    print("Content model file found:", content_model_path)
else:
    print("Content model file not found!")

if os.path.exists(cf_predictions_path):
    print("CF predictions file found:", cf_predictions_path)
else:
    print("CF predictions file not found!")
```

Content model file found: /content/drive/MyDrive/movie_dataset_public_final/xgb_cb_model.pkl
CF predictions file found: /content/drive/MyDrive/movie_dataset_public_final/cleaned_ratings_cosine_train.pkl

```
In [ ]: import joblib

content_model = joblib.load(content_model_path)

cf_predictions = pd.read_pickle(cf_predictions_path)
```

```
In [ ]: import pandas as pd

train_data_path = '/content/drive/MyDrive/movie_dataset_public_final/filt
train_data = pd.read_csv(train_data_path)
```

```
In [ ]: train_data
```

Out []:

	item_id	user_id	rating
0	5	997206	3.0
1	10	997206	4.0
2	13	997206	4.0
3	17	997206	5.0
4	21	997206	4.0
...
23514990	97938	187144	3.0
23514991	98809	187144	3.0
23514992	99114	187144	3.0
23514993	102445	187144	4.0
23514994	104841	187144	3.0

23514995 rows × 3 columns

```
In [ ]: train_data_sample = train_data.sample(n=10000, random_state=42)
train_data_sample
```

Out []:

	item_id	user_id	rating
9507762	3996	306188	3.0
4777740	2827	939224	3.0
1982187	592	41489	3.0
4015247	930	347251	5.0
21325961	48394	552153	5.0
...
7099488	337	702670	4.0
23131247	3753	242955	3.0
16881896	7386	48864	4.5
19204730	91500	906673	3.0
13463269	2124	138135	4.0

10000 rows × 3 columns

```
In [ ]: base_dir = "/content/drive/MyDrive/movie_dataset_public_final/raw/"

def load_ndjson_file(file_name, limit=None):
    data = []
    file_path = os.path.join(base_dir, file_name)

    with open(file_path, 'r', encoding='utf-8') as f:
        for i, line in enumerate(f):
```

```

        if limit and i >= limit:
            break
        try:
            data.append(json.loads(line))
        except json.JSONDecodeError as e:
            print(f"Error parsing line {i}: {line}")
            print(f"Error: {e}")
    return data

metadata = load_ndjson_file("metadata.json")

```

```

In [ ]: metadata_updated = load_ndjson_file("metadata_updated.json")
        metadata = pd.DataFrame(metadata_updated)
        print(metadata.head())

```

	title	directedBy	\
0	Toy Story (1995)	John Lasseter	
1	Jumanji (1995)	Joe Johnston	
2	Grumpier Old Men (1995)	Howard Deutch	
3	Waiting to Exhale (1995)	Forest Whitaker	
4	Father of the Bride Part II (1995)	Charles Shyer	

	starring	avgRating	imdbId
\			
0	Tim Allen, Tom Hanks, Don Rickles, Jim Varney,...	3.89146	0114709
1	Jonathan Hyde, Bradley Pierce, Robin Williams,...	3.26605	0113497
2	Jack Lemmon, Walter Matthau, Ann-Margret , Sop...	3.17146	0113228
3	Angela Bassett, Loretta Devine, Whitney Housto...	2.86824	0114885
4	Steve Martin, Martin Short, Diane Keaton, Kimb...	3.07620	0113041

	item_id
0	1
1	2
2	3
3	4
4	5

```

In [ ]: movie_features_path = '/content/drive/MyDrive/movie_dataset_public_final/
        movie_features = pd.read_csv(movie_features_path)

```

```

In [ ]: movie_features
        movie_features= movie_features.drop(columns=['avgRating', 'imdbId'])

```

```

In [ ]: movie_features

```

Out []:

	item_id	avg_negative	avg_neutral	avg_positive	avg_compound	directedB
0	1	-1.747673	0.447609	1.122360	1.293096	
1	2	-0.797492	-0.071923	0.741227	0.891817	
2	3	-0.701273	-0.143832	0.722673	1.028139	
3	4	-0.082240	0.143450	-0.056040	-0.207967	
4	5	-1.435475	-0.364290	1.530462	0.990010	
...	
3798	143385	-0.319489	1.003223	-0.562908	0.343855	
3799	148626	0.477404	0.636112	-0.937548	-0.406321	
3800	164179	-0.179858	1.577683	-1.161556	-0.234032	
3801	166528	1.903981	-1.282985	-0.560504	-1.134761	
3802	166643	-1.066447	1.285089	-0.157634	0.866084	

3803 rows × 318 columns

In []: `trained_feature_names = content_model.feature_names_in_`In []: `trained_feature_names`


```

Out[ ]: array(['user_id', 'avg_negative', 'avg_neutral', 'avg_positive',
               'avg_compound', 'directedBy_Alfred_Hitchcock',
               'directedBy_Barry_Levinson', 'directedBy_Clint_Eastwood',
               'directedBy_Joel_Schumacher', 'directedBy_Martin_Scorsese',
               'directedBy_Mike_Nichols', 'directedBy_Norman_Jewison',
               'directedBy_Oliver_Stone', 'directedBy_Other',
               'directedBy_Richard_Donner', 'directedBy_Ridley_Scott',
               'directedBy_Rob_Reiner', 'directedBy_Robert_Altman',
               'directedBy_Robert_Stevenson', 'directedBy_Ron_Howard',
               'directedBy_Sidney_Lumet', 'directedBy_Spike_Lee',
               'directedBy_Stanley_Kubrick', 'directedBy_Steven_Soderbergh',
               'directedBy_Steven_Spielberg', 'directedBy_Woody_Allen', '',
               'Bill_Murray', 'Bruce_Willis', 'Christopher_Walken',
               'Donald_Sutherland', 'Dustin_Hoffman', 'Gene_Hackman',
               'Harvey_Keitel', 'Jack_Nicholson', 'Julianne_Moore',
               'Meryl_Streep', 'Michael_Caine', 'Morgan_Freeman', 'Other',
               'Philip_Seymour_Hoffman', 'Robert_De_Niro', 'Robert_Duvall',
               'Robin_Williams', 'Samuel_L_Jackson', 'Sean_Connery', 'Tom_Hank
s',
               '1980s', '70mm', 'action', 'adapted_from_book', 'adultery',
               'adventure', 'atmospheric', 'based_on_a_book',
               'based_on_a_true_story', 'betrayal', 'black_and_white',
               'black_comedy', 'blood', 'boring', 'chase', 'children',
               'cinematography', 'classic', 'comedy', 'coming_of_age',
               'corruption', 'crime', 'criterion', 'cult_film', 'dancing', 'dar
k',
               'dark_comedy', 'death', 'dialogue', 'directorial_debut', 'disne
y',
               'dog', 'drama', 'drugs', 'england', 'family', 'fantasy',
               'father_daughter_relationship', 'father_son_relationship',
               'franchise', 'friendship', 'fun', 'funny', 'great_acting',
               'great_soundtrack', 'high_school', 'hilarious', 'history',
               'horror', 'humorous', 'imdb_top_250', 'independent_film',
               'infidelity', 'investigation', 'long', 'los_angeles', 'love',
               'marriage', 'money', 'murder', 'music', 'musical', 'mystery',
               'new_york', 'new_york_city', 'nudity', 'nudity_topless_brief_',
               'nudity_topless_', 'overrated', 'period_piece', 'police',
               'politics', 'predictable', 'prison', 'quirky', 'revenge',
               'romance', 'satire', 'sci-fi', 'sequel', 'serial_killer', 'sex',
               'sexuality', 'silly', 'slow', 'small_town', 'social_commentary',
               'story', 'stupid', 'stylized', 'suicide', 'surreal', 'suspense',
               'tense', 'thriller', 'true_story', 'violence',
               'visually_appealing', 'war', 'witty', 'about', 'adventures', 'al
l',
               'america', 'american', 'an', 'and', 'at', 'baby', 'bad', 'big',
               'black', 'blue', 'boys', 'bride', 'by', 'cat', 'city', 'day',
               'days', 'de', 'dead', 'der', 'die', 'down', 'fire', 'first', 'fo
r',
               'friday', 'from', 'girl', 'good', 'great', 'harry', 'heart',
               'high', 'home', 'house', 'ii', 'iii', 'in', 'is', 'it', 'kid',
               'king', 'kiss', 'la', 'last', 'le', 'legend', 'les', 'life',
               'little', 'lost', 'mad', 'man', 'me', 'men', 'movie', 'mr', 'mr
s',
               'my', 'new', 'night', 'no', 'of', 'on', 'one', 'out', 'part',
               'planet', 'red', 'return', 'road', 'secret', 'star', 'street',
               'that', 'the', 'thing', 'three', 'time', 'to', 'trek', 'two', 'u
p',
               'what', 'white', 'who', 'wild', 'with', 'world', 'you',
               'release_decade_1920', 'release_decade_1930',
               'release_decade_1940', 'release_decade_1950',

```

```
'release_decade_1960', 'release_decade_1970',
'release_decade_1980', 'release_decade_1990',
'release_decade_2000', 'release_decade_2010'], dtype='<U28')
```

```
In [ ]: def get_collab_prediction(user_id, item_id, cf_data):
        """
        Retrieve the collaborative filtering prediction for a given user and
        """
        try:
            prediction = cf_data.loc[(cf_data['user_id'] == user_id) & (cf_da
        except IndexError:
            prediction = cf_data['restored_predicted_rating'].mean()
        return prediction
```

```
In [ ]: trained_feature_names = content_model.feature_names_in_
print("Feature names in trained model:", trained_feature_names)
print("Columns in movie_features:", movie_features.columns)

valid_trained_feature_names = [name for name in trained_feature_names if

missing_columns = [col for col in trained_feature_names if col not in mov
for col in missing_columns:
    movie_features[col] = 0

filtered_movie_features = movie_features[valid_trained_feature_names]
print("Filtered movie_features columns:", filtered_movie_features.columns)
```

Feature names in trained model: ['user_id' 'avg_negative' 'avg_neutral' 'avg_positive' 'avg_compound'
'directedBy_Alfred_Hitchcock' 'directedBy_Barry_Levinson'
'directedBy_Clint_Eastwood' 'directedBy_Joel_Schumacher'
'directedBy_Martin_Scorsese' 'directedBy_Mike_Nichols'
'directedBy_Norman_Jewison' 'directedBy_Oliver_Stone' 'directedBy_Other'
'directedBy_Richard_Donner' 'directedBy_Ridley_Scott'
'directedBy_Rob_Reiner' 'directedBy_Robert_Altman'
'directedBy_Robert_Stevenson' 'directedBy_Ron_Howard'
'directedBy_Sidney_Lumet' 'directedBy_Spike_Lee'
'directedBy_Stanley_Kubrick' 'directedBy_Steven_Soderbergh'
'directedBy_Steven_Spielberg' 'directedBy_Woody_Allen' '' 'Bill_Murray'
'Bruce_Willis' 'Christopher_Walken' 'Donald_Sutherland' 'Dustin_Hoffman'
'Gene_Hackman' 'Harvey_Keitel' 'Jack_Nicholson' 'Julianne_Moore'
'Meryl_Streep' 'Michael_Caine' 'Morgan_Freeman' 'Other'
'Philip_Seymour_Hoffman' 'Robert_De_Niro' 'Robert_Duvall'
'Robin_Williams' 'Samuel_L_Jackson' 'Sean_Connery' 'Tom_Hanks' '1980s'
'70mm' 'action' 'adapted_from_book' 'adultery' 'adventure' 'atmospheric'
'based_on_a_book' 'based_on_a_true_story' 'betrayal' 'black_and_white'
'black_comedy' 'blood' 'boring' 'chase' 'children' 'cinematography'
'classic' 'comedy' 'coming_of_age' 'corruption' 'crime' 'criterion'
'cult_film' 'dancing' 'dark' 'dark_comedy' 'death' 'dialogue'
'directorial_debut' 'disney' 'dog' 'drama' 'drugs' 'england' 'family'
'fantasy' 'father_daughter_relationship' 'father_son_relationship'
'franchise' 'friendship' 'fun' 'funny' 'great_acting' 'great_soundtrack'
'high_school' 'hilarious' 'history' 'horror' 'humorous' 'imdb_top_250'
'independent_film' 'infidelity' 'investigation' 'long' 'los_angeles'
'love' 'marriage' 'money' 'murder' 'music' 'musical' 'mystery' 'new_york'
'new_york_city' 'nudity' 'nudity_topless_brief_' 'nudity_topless_'
'overrated' 'period_piece' 'police' 'politics' 'predictable' 'prison'
'quirky' 'revenge' 'romance' 'satire' 'sci-fi' 'sequel' 'serial_killer'
'sex' 'sexuality' 'silly' 'slow' 'small_town' 'social_commentary' 'story'
'stupid' 'stylized' 'suicide' 'surreal' 'suspense' 'tense' 'thriller'
'true_story' 'violence' 'visually_appealing' 'war' 'witty' 'about'
'adventures' 'all' 'america' 'american' 'an' 'and' 'at' 'baby' 'bad'
'big' 'black' 'blue' 'boys' 'bride' 'by' 'cat' 'city' 'day' 'days' 'de'
'dead' 'der' 'die' 'down' 'fire' 'first' 'for' 'friday' 'from' 'girl'
'good' 'great' 'harry' 'heart' 'high' 'home' 'house' 'ii' 'iii' 'in' 'is'
'it' 'kid' 'king' 'kiss' 'la' 'last' 'le' 'legend' 'les' 'life' 'little'
'lost' 'mad' 'man' 'me' 'men' 'movie' 'mr' 'mrs' 'my' 'new' 'night' 'no'
'of' 'on' 'one' 'out' 'part' 'planet' 'red' 'return' 'road' 'secret'
'star' 'street' 'that' 'the' 'thing' 'three' 'time' 'to' 'trek' 'two'
'up' 'what' 'white' 'who' 'wild' 'with' 'world' 'you'
'release_decade_1920' 'release_decade_1930' 'release_decade_1940'
'release_decade_1950' 'release_decade_1960' 'release_decade_1970'
'release_decade_1980' 'release_decade_1990' 'release_decade_2000'
'release_decade_2010']

Columns in movie_features: Index(['item_id', 'avg_negative', 'avg_neutral', 'avg_positive',
'avg_compound', 'directedBy_Alan_Parker', 'directedBy_Alfred_Hitchcock',
'directedBy_Ang_Lee', 'directedBy_Barry_Levinson',
'directedBy_Billy_Wilder',
...,
'release_decade_1920', 'release_decade_1930', 'release_decade_1940',
'release_decade_1950', 'release_decade_1960', 'release_decade_1970',
'release_decade_1980', 'release_decade_1990', 'release_decade_2000',
'release_decade_2010']

```

        'release_decade_2010'],
        dtype='object', length=318)
Filtered movie_features columns: Index(['avg_negative', 'avg_neutral', 'avg_
g_positive', 'avg_compound',
        'directedBy_Other', 'Other', '1980s', '70mm', 'action', 'adultery',
        ...
        'release_decade_1920', 'release_decade_1930', 'release_decade_194
0',
        'release_decade_1950', 'release_decade_1960', 'release_decade_197
0',
        'release_decade_1980', 'release_decade_1990', 'release_decade_200
0',
        'release_decade_2010'],
        dtype='object', length=181)

```

```

In [ ]: def get_content_prediction(item_id, movie_features, content_model):
        """
        Predict the average rating of a movie using the content-based model.
        """

        features = movie_features.loc[movie_features['item_id'] == item_id].d

        features = features[content_model.feature_names_in_]
        return content_model.predict(features)[0]

```

```

In [ ]: def predict_hybrid_score(user_id, item_id, cf_data, movie_features, conte
        """
        Combine predictions from content-based and collaborative filtering mo
        """

        content_pred = get_content_prediction(item_id, movie_features, conten

        collab_pred = get_collab_prediction(user_id, item_id, cf_data)

        beta = 1 - alpha
        return alpha * content_pred + beta * collab_pred

```

```

In [ ]: filtered_train_data = train_data_sample[train_data_sample['item_id'].isin

        print(f"Original train_data size: {len(train_data_sample)}")
        print(f"Filtered train_data size: {len(filtered_train_data)}")

```

Original train_data size: 10000
 Filtered train_data size: 9975

```

In [ ]: from sklearn.metrics import mean_squared_error
        import numpy as np

        from sklearn.metrics import mean_squared_error
        import numpy as np

        def compute_mse_rmse(alpha, test_data, cf_data, movie_features, content_m
        """
        Compute MSE and RMSE for test samples, considering two cases:
        1. Samples with only CB prediction.
        2. Samples with both CB and CF predictions.
        """

        true_ratings_cb_only = []
        cb_only_predictions = []
        true_ratings_hybrid = []
        hybrid_predictions = []

```

```

for _, row in test_data.iterrows():
    user_id = row['user_id']
    item_id = row['item_id']
    true_rating = row['rating']

    content_pred = get_content_prediction(item_id, movie_features, co

    try:
        collab_pred = get_collab_prediction(user_id, item_id, cf_data
    except KeyError:
        collab_pred = None

    if collab_pred is None:

        true_ratings_cb_only.append(true_rating)
        cb_only_predictions.append(content_pred)
    else:

        hybrid_pred = alpha * content_pred + (1 - alpha) * collab_pre
        true_ratings_hybrid.append(true_rating)
        hybrid_predictions.append(hybrid_pred)

mse_cb_only = mean_squared_error(true_ratings_cb_only, cb_only_predic
mse_hybrid = mean_squared_error(true_ratings_hybrid, hybrid_predictio

total_samples = len(test_data)
mse = (mse_cb_only * len(true_ratings_cb_only) + mse_hybrid * len(tru
rmse = np.sqrt(mse)
return mse, rmse

```

In []: `from tqdm import tqdm`

```

sample_train_data = train_data

alphas = np.linspace(0.45, 0.6, 8)
results = []

for alpha in alphas:

    mse, rmse = compute_mse_rmse(alpha, sample_train_data, cf_predictions
    results.append((alpha, mse, rmse))
    print(f"Alpha: {alpha:.2f}, MSE: {mse:.4f}, RMSE: {rmse:.4f}")

best_alpha, best_mse, best_rmse = min(results, key=lambda x: x[2])
print(f"Best alpha: {best_alpha:.2f}, Best MSE: {best_mse:.4f}, Best RMSE

```

```
/usr/local/lib/python3.10/dist-packages/xgboost/core.py:158: UserWarning:  
[15:44:34] WARNING: /workspace/src/common/error_msg.cc:58: Falling back to  
prediction using DMatrix due to mismatched devices. This might lead to hig  
her memory usage and slower performance. XGBoost is running on: cuda:0, wh  
ile the input data is on: cpu.
```

Potential solutions:

- Use a data structure that matches the device ordinal in the booster.
- Set the device for booster before call to inplace_predict.

This warning will only be shown once.

```
warnings.warn(smsg, UserWarning)
```

Alpha: 0.45, MSE: 0.9525, RMSE: 0.9759

In []: movie_features

```
In [29]: import joblib
import json
import os

# Define paths to save files
output_dir = "/content/drive/MyDrive/movie_dataset_public_final/"
os.makedirs(output_dir, exist_ok=True) # Create directory if it doesn't

# Save the content-based model
content_model_path = os.path.join(output_dir, "hybridwithalpha_model.pkl")
joblib.dump(content_model, content_model_path)
print(f"Content model saved at: {content_model_path}")

# Save best alpha and evaluation metrics
results_path = os.path.join(output_dir, "results.json")
results_data = {
    "best_alpha": best_alpha,
    "best_mse": best_mse,
    "best_rmse": best_rmse,
    "all_results": results
}
with open(results_path, "w") as f:
    json.dump(results_data, f, indent=4)
print(f"Results saved at: {results_path}")
```

Content model saved at: /content/drive/MyDrive/movie_dataset_public_final/hybridwithalpha_model.pkl

Results saved at: /content/drive/MyDrive/movie_dataset_public_final/results.json

```
In [32]: def add_predictions(data, alphas, cf_data, movie_features, content_model):
data = data.copy()
for alpha in alphas:
    col_name = f"predicted_score_alpha_{alpha:.2f}"
    data[col_name] = data.apply(
        lambda row: predict_hybrid_score(
            row['user_id'], row['item_id'], cf_data, movie_features,
        ),
```

```

        axis=1
    )
    return data

predicted_data = add_predictions(sample_train_data, [best_alpha], cf_pred
predicted_data_path = os.path.join(output_dir, "predicted_train_data.csv")
predicted_data.to_csv(predicted_data_path, index=False)
print(f"Predicted data saved at: {predicted_data_path}")

```

Predicted data saved at: /content/drive/MyDrive/movie_dataset_public_final/predicted_train_data.csv

```

In [36]: def generate_recommendations(user_id, movie_id, cf_data, movie_features,
    """
    Generate top-N movie recommendations for a user, considering a specific movie.
    """
    recommendations = []
    user Rated items = cf_data[cf_data['user_id'] == user_id]['item_id'].

    for item_id in movie_features['item_id'].unique():
        if item_id not in user Rated items and item_id != movie_id: # Example
            # Predict the hybrid score for this user-item pair
            score = predict_hybrid_score(user_id=user_id, item_id=item_id,
                                         cf_data=cf_data,
                                         movie_features=movie_features,
                                         content_model=content_model,
                                         alpha=alpha)
            recommendations.append((item_id, score))

    # Sort recommendations by score in descending order
    recommendations = sorted(recommendations, key=lambda x: x[1], reverse=True)

    return recommendations

# Example usage
user_id = 19 # Replace with a valid user_id from your dataset
movie_id = 1276 # Replace with a valid movie_id the user is interested in
top_recommendations = generate_recommendations(user_id, movie_id, cf_predicted_data, movie_features)
print("Top Recommendations for User, considering movie {}: ".format(movie_id))

```

Top Recommendations for User, considering movie 1276: [(26131, 4.132750265217117)]

```

In [38]: def generate_recommendations3(user_id, movie_id, cf_data, movie_features,
    """
    Generate top-N movie recommendations for a user, considering a specific movie.
    """
    recommendations = []
    user Rated items = cf_data[cf_data['user_id'] == user_id]['item_id'].

    for item_id in movie_features['item_id'].unique():
        if item_id not in user Rated items and item_id != movie_id:
            score = predict_hybrid_score(user_id=user_id, item_id=item_id,
                                         cf_data=cf_data,
                                         movie_features=movie_features,
                                         content_model=content_model,
                                         alpha=alpha)
            recommendations.append((item_id, score))

    recommendations = sorted(recommendations, key=lambda x: x[1], reverse=True)

```



```
    return recommendations

user_id = 19
movie_id = 1276
top_recommendations = generate_recommendations3(user_id, movie_id, cf_pre
print("Top Recommendations for User, considering movie {}".format(movie_
```

Top Recommendations for User, considering movie 1276: [(26131, 4.132750265217117), (1148, 4.044837548351577), (2932, 3.931541263675979)]

In []: